

DSCI-633 Foundations of Data Science & Analytics Project

December 2021

Title:

**ANALYSING WATCH TIMES OF HOTSTAR USERS AND
PREDICTING THE USER SENTIMENT**

Team Members:

**Name: Omkar Khanvilkar
Email: osk8557@rit.edu
RIT ID: 347001066**

**Name: Sujan Dutta
Email: sd2516@g.rit.edu
RIT ID: 782000242**

**Name: Pranav Nair
Email: pn1922@g.rit.edu
RIT ID: 359006909**

**Name: Varun Tandon
Email: vt9438@g.rit.edu
RIT ID: 572002978**

DEVELOPMENT OF QUESTION/HYPOTHESIS

Over the past decade, we have seen an increase in viewership for streaming services across the globe. Sites like Netflix, Hulu, and Hotstar dominate the streaming market with the content that they provide. The common denominator across these sites is the userbase and the feedback they provide to get a better experience.

Hotstar is a platform that has over 100 million user count, and more than 35,000 hours of content accessible over a variety of genres. As is the case with most major streaming services, using the sheer scale of users and content to create tailor-made recommendations and content for users generates a lot of value for Hotstar. Thus comes our problem statement.

How does a service like Hotstar capture its users' sentiments? Regular questioning often breaks immersion and can cause a biased answer. To approach this problem, we look at an alternate method, deducing the user sentiment by analysing the user behaviour.

The data the user generates is watch time in seconds, over multiple parameters. The users were also asked if they had a positive or a negative sentiment based on their overall experience. Our goal is to find the correlation between these parameters and the sentiment.

DATA RESEARCH

Upon searching for relevant data for the proposed problem, we found a dataset on HackerEarth that comprises of two files:

- *train_df.json*, which consists of watch times and the sentiment for 200,000 users.
- *test_df.json*, which consists of the sentiments for 100,000 users.

The data itself consists of 7 unique columns:

- *id*: A unique identifier variable.
- *titles*: Titles of the show watched by the user along with the watch time.
- *genres*: Names of the genres watched by the user along with the watch time.
- *cities*: Names of the cities the user was present in along with the watch time.
- *tod*: The time of the day in hours the user began watching along with the watch time.
- *dow*: The day of the week the user began watching along with the watch time.
- *segment*: The sentiment of the user (as positive or negative.)

The data itself is open to the public, present on HackerEarth.

LITERATURE REVIEW

The topic of customer (user) classification and segmentation has gained a lot of research attention in the last two decades because of the emergence of multiple OTT platforms. It all started with the release of the famous Netflix dataset [1] that contains 100M movie ratings in 2006. Since then, many researchers have used this dataset to build movie recommendation systems and customer segmentation models. However, our dataset is a much newer one and it comes from Hotstar, a different OTT platform. While doing literature survey we found that there exist many works that use techniques like matrix factorization, nearest neighbors' approach [2], clustering approaches [3] to recommend movies (or products), and perform customer segmentation. In this work, however, we are treating the problem as a classification problem as the data contains a label denoting the segment (positive or negative) corresponding to a customer. Treating the problem as a classification task will allow us to use the popular tree-based algorithms like Random Forest and XGBoost.

ANALYSIS: STRATEGY

There were two main issues faced with this data set. Primarily, the data was purely categorical in nature, in the form of "*value:watchtime*", separated by commas. Preprocessing this kind of data tends to be pretty time consuming and heavy on the computational load.

To resolve this, the preprocessing was done on the json files itself, that allowed the local systems to utilize the most of their computational power without wasting computational time.

The preprocessing itself was to create a function that parsed through the row of a single user, separate the value and the watch time, and add the value as a column with the user as the row, and the watch time as the data for that cell. With each user parsed, common columns were merged together, and a sparse data frame was created, with every single value type as a new column.

The second issue faced was the heavy class imbalance. The negative sentiments overshadowed the positive sentiments with a 23:2 ratio. With such a heavy imbalance, the models were sure to overfit and predict the incorrect class.

The solution proposed for this issue was to oversample the positive class and undersample the negative class. The oversampling technique used here was SMOTE. After doing so, the ratio of the negative class to the positive class reduced to 5:2.

The main features used for the final dataset were:

- cities, genres, titles, dow, tod

Upon further analysis, the following assumptions were made based on EDA:

- The cities with users of positive sentiment did not follow the trend of the cities with users of negative sentiment. Furthermore, there were more cities with negative sentiment than positive.
- Cricket was the highest watched genre, but there was no other sport in the top ten most-watched genres.
- The most-watched hour of the day was 9pm, followed by 10pm
- The most-watched day of the week was Saturday, followed by Friday. This indicates that the users tend to watch on weekends rather than weekdays. This also corroborates

with the time of the day analysis; Users watch on weekend nights, but not on Sunday nights.

ANALYSIS: CODE

After pre-processing, the final number of columns ended up being 1448. To reduce this number, feature engineering was performed on the dataset.

Each feature was approached in order with the following results:

- Genres:
 - All sports (except for cricket) were combined into a single column, with the watch times being summed.
 - A *genre_count* feature was added to keep track of the total genres watched by the user.
- Cities:
 - Only the top 15 cities based on watch time were added to the final dataset.
- dow:
 - Two new columns were added: *total_watch_time*, to keep track of the user's total watch time of the week, and *count*, to keep track of the number of days the user has watched for.
- tod:
 - Only the top 5 hours based on watch time were added to the final dataset
 - An additional column for the total hours watched by the user.
- titles:
 - An additional column for the total number of titles watched by the user

The final data frame consisted of 78 unique features excluding the target after the engineering. The over and under sampling was done to this data frame.

Next, the four models were tested on this data frame. Each model was trained without any tuning and then validated on a validation data set. Then the model underwent hyperparameter tuning and was validated on a validation data set.

These were the final results after the tuning:

- Logistic Regression:
 - Without Hyperparameter tuning: 0.78020
 - With Hyperparameter tuning: 0.78029
- Decision Tree:
 - Without Hyperparameter tuning: 0.6204
 - With Hyperparameter tuning: 0.6682
- Random Forest:
 - Without Hyperparameter tuning: 0.8068
 - With Hyperparameter tuning: 0.8056

- XGBoost:
 - Without Hyperparameter tuning: 0.80
 - With Hyperparameter tuning: 0.8119

All results above are indicative of the AUC ROC score.

Thus, XGBoost was the model chosen to predict on our test data as it had the highest ROC AUC score amongst the four.

The next step was to deploy the model, using a python library, “*streamlit*”. This is an open-source framework used for building web apps for machine learning and data science models.

The UI of the web app allows input of the user number, and can predict the sentiment of the user.

WORK PLANNING AND ORGANIZATION

The entire workload was split equally amongst the four team members. Each part of the project was compartmentalized and broken into 4 equal parts, i.e., the pre-processing, EDA and model engineering was tackled by all 4 members.

IMPROVING TEAMWORK AND COLLABORATION

As mentioned above, the workload for each phase of the project was divided equally. For the preprocessing phase, we came up with the concept of parsing the first 100 rows of the data frame, and then decided to preprocess on the json file itself.

For the EDA, I worked on the tod feature, i.e., the time of the day feature. I analysed what trends were prominent in the dataset and how the watch time affected both the other features as well as the segment.

For the training, I focused on the Random Forest Classifier. I trained a model with default parameters to establish a baseline and then performed hyperparameter tuning on the RF model. The tuning however resulted in a 0.0012% loss, which was chalked up to error.

Overall, there was equal contribution for each segment, including the presentation, and the final notebook itself. All four members of the group worked on collaborating on the notebook with their individual findings.

INDIVIDUAL CONTRIBUTION

Part 1 : Feature Engineering.

The data was extracted and the data in titles were used as features in the title column. But, unlike the other columns, I couldn't do it straight since taking all of the titles as features resulted in almost 2000 features. It boosted a dataset's sparsity and complexity. To address the above mentioned issue, I opted to only include titles that had the majority of user watch time. To do so, I computed watch time for all of the titles and discovered that the bulk of cricket match titles received a high user watch time. This finding was evident since Hotstar broadcasts cricket matches in India, and cricket is the most popular sport in the country. As a result, I regarded cricket as a whole new feature in the genres section for all the cricket watch time. Along with all this changes I've also included a column that shows the total number of titles the user has seen. So that we can quickly determine how many type of content the user is seeing.

The DOW column in the train dataset indicates the user's day of week watch time. I've added two new columns to DOW for the feature engineering component, which indicate the user's overall watch time, which is calculated by adding all of the weeks' watch time. We can determine how many people are active on the platform and whether or not they are active users, allowing us to construct a list of active users and extract more information from them.

Part 2 : EDA.

In the EDA section, I contributed to the genres section, where I determined that cricket is the most popular genre among Hotstar viewers. I created a graph that shows the top 10 genres based on the users' viewing time.

Part 4 : Getting the instance to train model

We attempted to train the model on a local machine, however because to the large amount of data, the operation was always terminated with an error. So I've finally come up with a GPU instance that will assist us all in training our models.

Part 3 : Model Training.

I trained four models: logistic regression with and without hyper parameter tuning, decision tree with and without hyper parameter tuning, and decision tree with and without hyper parameter tuning.

Logistic Regression :

The goal behind logistic regression is to discover the likelihood in all of the features we come up with during data pre-processing, and then compare it to the probability of certain outcomes. I started with a standard logistic regression model, which yielded a score of 0.7802. I hypertunned it even further to achieve a better score, then I used the sklearn GridSearchCV package to determine the optimum parameter for hypertunning. Where I got the score of 0.78029

Decision Tree:

Decision tree are used in which previous outcomes are used for the next decision to make so basically it's decision tree. I used the conventional decision tree technique to train the model, which yielded a score of 0.6204. I attempted hyperparameter tuning in the decision tree to enhance the score. I utilized the sklearn GridSearchCV library function in the decision tree. This works by supplying the default hyperparameter, which is then gives with the optimal parameter that yields the highest model training score. I utilized the best parameters and hypertuned the model training to get a 0.6682 score.

REFERENCE

- [1] J. Bennett, S. Lanning, et al., "The netflix prize," in Proceedings of KDD cup and workshop, vol. 2007, p. 35, New York, NY, USA., 2007.
- [2] G. Takács, I. Pilászy, B. Németh, and D. Tikk, "Matrix factorization and neighbor based algorithms for the netflix prize problem," in Proceedings of the 2008 ACM conference on Recommender systems, pp. 267–274, 2008.
- [3] P. Q. Brito, C. Soares, S. Almeida, A. Monte, and M. Byvoet, "Customer segmentation in a large database of an online customized fashion business," Robotics and Computer-Integrated Manufacturing, vol. 36, pp. 93–100, 2015